


GRAPHCOMPOSE 2.0 / MODULE-FIRST

Compose once. Render through modules.

A lean Java document core with a semantic model, deterministic layout geometry and opt-in output backends.

THE 2.0 MODULE GRAPH



io.github.demchaav:graph-compose-core

+ choose a renderer module

● graph-compose-core
lean engine


● graph-compose
drop-in PDF

● ServiceLoader SPI
bring a renderer

● layoutSnapshot()
geometry regression

Semantic at the core. Format at the edge.

The canonical DSL produces one semantic DocumentGraph. Fixed-layout output resolves geometry; semantic exporters consume the graph directly.



The renderer boundary is explicit; Android or another backend is an extension opportunity, not a shipped 2.0 claim.

Choose only what you need.

Eight engine-train artifacts move in lockstep. Fonts and emoji keep independent version lines; heavy render and template dependencies are opt-in.

DROP-IN PDF DEFAULT

graph-compose

Core + PDF renderer. Canonical callers keeping this coordinate retain PDF out of the box.

LEAN / CUSTOM LEAN

graph-compose-core

Authoring, semantic model and layout services. Add a renderer module only when needed.

BATTERIES INCLUDED BUNDLE

graph-compose-bundle

Default PDF stack plus templates, bundled fonts and emoji. Office exporters stay opt-in.

THE LOCKSTEP 2.0 TRAIN

One version, smaller dependency trees, explicit format boundaries.

● core

authoring + model

● render-pdf

fixed PDF

● render-docx

semantic DOCX

● render-pptx

manifest only

● templates

opt-in presets

● testing

snapshot tools

● graph-compose

PDF wrapper

● bundle

PDF + templates + assets

Accepted migration trade-off: built-in templates are no longer carried by graph-compose; add graph-compose-templates or use the bundle.

2 / 4

Measured locally. Read with context.

Equivalent 1000-row report fixtures from the bundled 2026-06-15 single-machine snapshot. These values are directional, not a final 2.0 marketing benchmark.

45.3 ms

GraphCompose avg render
1000 rows

15.7 MiB



avg allocated / render
thread allocation

40 / 200 / 1000

report rows
scaling fixtures

2026-06-15

snapshot date
provenance incomplete





Scope: one developer machine; CPU, OS, JDK and commit were not stored in this historical file. The JSON stores 50/100 only for its small-invoice tier and omits scaling iteration metadata. Refresh both data and provenance on the final 2.0 commit before release claims.

Reproduce: scripts/run-benchmarks.ps1 - Current docs: docs/operations/benchmarks.md

Watch the curve, not one headline number.

The same local reference run across 40, 200 and 1000 rows. Allocation values are independent per-render measurements; they are not shares of a common heap.



5.38x
iText 9 took this many times longer in the 1000-row fixture

4.19x
iText 9 allocated this many times more in the same fixture

6.6%
latency gap vs JasperReports - inside the documented 5-10% local noise band

What 2.0 can claim today: **modular packaging, renderer isolation and deterministic layout snapshots**. Refresh performance evidence separately on the release commit.

GraphCompose renders this deck itself: gradients, canvas layers, SVG, clipping and native charts.